





Welcome to

1. Integration intro, Java Apps, Tomcat, XML config

KEA System Integration F2020 10 ECTS

Henrik Kramselund Jereminsen hkj@zencurity.com @kramse  

Slides are available as PDF, [kramse@Github](https://github.com/kramse)

1-Introduction-system-integration.tex in the repo security-courses

Plan for today




- Integration intro
- TCP/IP, DNS, HTTP intro
- Java Apps, Tomcat, XML config
- Git intro

Exercises

- TCP/IP
- Apache Tomcat
- Git introduction

Reading Summary



Enterprise Integration Patterns, Gregor Hohpe and Bobby Woolf, 2004

ISBN: 978-0-321-20068-6 EIP for short

EIP book chapter 1-2

Definition: System integration



System integration is defined in engineering as the process of bringing together the component sub-systems into one system (an aggregation of subsystems cooperating so that the system is able to deliver the overarching functionality) and ensuring that the subsystems function together as a system,[1] and in information technology[2] as the process of linking together different computing systems and software applications physically or functionally,[3] to act as a coordinated whole.

The system integrator integrates discrete systems utilizing a variety of techniques such as computer networking, enterprise application integration, business process management or manual programming.[4]

Source:

https://en.wikipedia.org/wiki/System_integration



"That's why Bobby Woolf and I documented a pattern language consisting of 65 integration patterns to establish a technology-independent vocabulary and a visual notation to design and document integration solutions. Each pattern not only presents a proven solution to a recurring problem, but also documents common "gotchas" and design considerations.

The patterns are brought to life with examples implemented in messaging technologies, such as JMS, SOAP, MSMQ, .NET, and other EAI Tools. The solutions are relevant for a wide range of integration tools and platforms, such as IBM WebSphere MQ, TIBCO, Vitria, WebMethods (Software AG), or Microsoft BizTalk, messaging systems, such as JMS, WCF, Rabbit MQ, or MSMQ, ESB's such as Apache Camel, Mule, WSO2, Oracle Service Bus, Open ESB, SonicMQ, Fiorano or Fuse ServiceMix."

Source:

<https://www.enterpriseintegrationpatterns.com/>

Integration intro



- System integration in information technology
- the process of linking together different computing systems and software applications
- Examples sales, inventory, procurement, human resources, email, web sites etc.

Why Enterprise Integration Patterns?



Enterprise integration is too complex to be solved with a simple 'cookbook' approach. Instead, patterns can provide guidance by documenting the kind of experience that usually lives only in architects' heads: they are accepted solutions to recurring problems within a given context. Patterns are abstract enough to apply to most integration technologies, but specific enough to provide hands-on guidance to designers and architects. Patterns also provide a vocabulary for developers to efficiently describe their solution.

Patterns are not 'invented'; they are harvested from repeated use in practice. If you have built integration solutions, it is likely that you have used some of these patterns, maybe in slight variations and maybe calling them by a different name. The purpose of this site is not to "invent" new approaches, but to present a coherent collection of relevant and proven patterns, which in total form an integration pattern language.

Source:

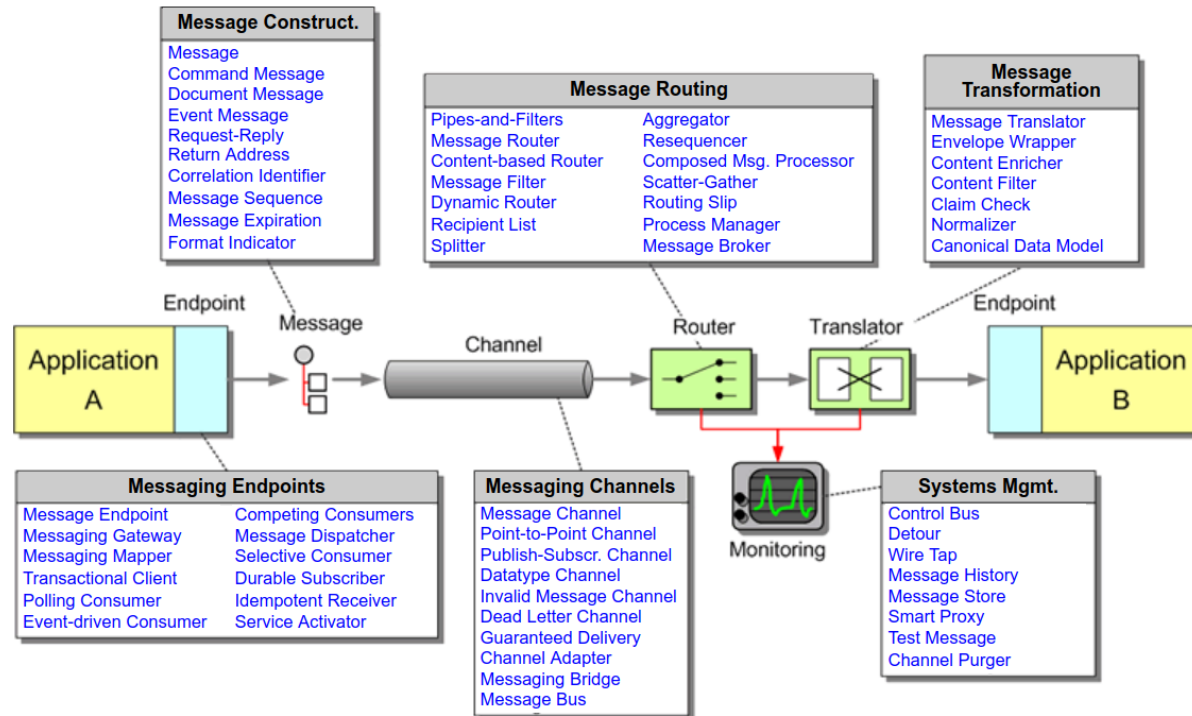
<https://www.enterpriseintegrationpatterns.com/>

EIP Patterns



Messaging Patterns

We have documented [65 messaging patterns](#), organized as follows:



Challenges



- Networks are unreliable. The internet is always broken, somewhere a link is down, a system being booted etc.
- Networks are slow. Sending data across networks are slower than making a local call
- Any two applications are different. Different programming languages, operating systems, and data formats
- Change is inevitable. Applications change over time
- Added: everything is linked, everything uses networking

Helpful patterns



- File Transfer(43)
- Shared database (47)
- Remote Procedure Invocation (50) - typically using Remote Procedure Call (RPC)
- Messaging (53) one application publishes a message to a common message channel, other applications read from the channel

Source: EIP book

Integration Styles



- Chapter 2 of the EIP book

Application coupling



Application coupling — Even integrated applications should minimize their dependencies on each other so that each can evolve without causing problems for the others. Tightly coupled applications make numerous assumptions about how the other applications work; when the applications change and break those assumptions, the integration breaks. The interface for integrating applications should be specific enough to implement useful functionality, but general enough to allow that implementation to change as needed.

Intrusivenss / Integration simplicity



Intrusivenss / Integration simplicity — When integrating an application into an enterprise, developers should strive to minimize changing the application and minimize the amount of integration code needed. Yet changes and new code will usually be necessary to provide good integration functionality, and the approaches with the least impact on the application may not provide the best integration into the enterprise.

Selecting Integration technology



Selecting Integration technology — Different integration techniques require varying amounts of specialized software and hardware. These special tools can be expensive, can lead to vendor lock-in, and increase the burden on developers to understand how to use the tools to integrate applications.

Data format



Data format — Integrated applications must agree on the format of the data they exchange, or must have an intermediate translator to unify applications that insist on different data formats. A related issue is data format evolution and extensibility—how the format can change over time and how that will affect the applications.

Data timeliness



Data timeliness — Integration should minimize the length of time between when one application decides to share some data and other applications have that data. Data should be exchanged frequently in small chunks, rather than waiting to exchange a large set of unrelated items. Applications should be informed as soon as shared data is ready for consumption. Latency in data sharing has to be factored into the integration design; the longer sharing can take, the more opportunity for shared data to become stale, and the more complex integration becomes.

Data or functionality




Data or functionality — Integrated applications may not want to simply share data, they may wish to share functionality such that each application can invoke the functionality in the others. Invoking functionality remotely can be difficult to achieve, and even though it may seem the same as invoking local functionality, it works quite differently, with significant consequences for how well the integration works.

Remote Communication / Asynchronicity



Remote Communication / Asynchronicity — Computer processing is typically synchronous, such that a procedure waits while its subprocedure executes. It's a given that the subprocedure is available when the procedure wants to invoke it. However, a procedure may not want to wait for the subprocedure to execute; it may want to invoke the subprocedure asynchronously, starting the subprocedure but then letting it execute in the background. This is especially true of integrated applications, where the remote application may not be running or the network may be unavailable—the source application may wish to simply make shared data available or log a request for a subprocedure call, but then go on to other work confident that the remote application will act sometime later.

File Transfer




File Transfer — Have each application produce files of shared data for others to consume, and consume files that others have produced.

Common systems and technologies used:

- File Transfer Protocol (FTP) - old protocol, uses clear text password - should not be used, but still is
- SFTP/SCP - replaces FTP, Secure FTP/ Secure Copy is part of the Secure Shell (SSH) protocol - available since 1995
- Hyper Text Transfer Protocol / HTTP Secure (HTTP/HTTPS) - web based protocols

Shared Database




Shared Database — Have the applications store the data they wish to share in a common database.

Common systems and technologies used:

- database management system (DBMS) using Structured Query Language (SQL), relational database examples:
- PostgreSQL, Oracle DM, Microsoft SQL, MySQL <https://en.wikipedia.org/wiki/SQL>
- NoSQL databases has been a new input with examples like: MongoDB, CouchDB, Redis, RIAK <https://en.wikipedia.org/wiki/NoSQL>


Remote Procedure Invocation



Remote Procedure Invocation — Have each application expose some of its procedures so that they can be invoked remotely, and have applications invoke those to run behavior and exchange data. Common systems and technologies used:

- Java remote method invocation (RMI), Unix RPC
- XMLHttpRequest (XHR) JavaScript in the browser makes connections and requests:
<https://en.wikipedia.org/wiki/XMLHttpRequest>
- Common Object Request Broker Architecture (CORBA) used in the 1990s but not very relevant anymore
- See more at https://en.wikipedia.org/wiki/Remote_procedure_call

Messaging



Messaging — Have each application connect to a common messaging system, and exchange data and invoke behavior using messages.


Common systems and technologies used:

- Java Message Service (JMS) API is a Java message-oriented middleware Application Programming Interface (API)
- Apache ActiveMQ, RabbitMQ, Oracle WebLogic
- See more at https://en.wikipedia.org/wiki/Message_passing

TCP/IP, DNS, HTTP intro




Networks



Our network is similar to regular networks, as found in enterprises

Internet Today



Clients and servers, roots in the academic world

Protocols are old, some more than 20 years

Very little is encrypted, mostly HTTPS

Internet is Open Standards!



We reject kings, presidents, and voting.
We believe in rough consensus and running code.
– The IETF credo Dave Clark, 1992.

Request for comments - RFC - er en serie af dokumenter

RFC, BCP, FYI, informational

de første stammer tilbage fra 1969

Ændres ikke, men får status Obsoleted når der udkommer en nyere version af en standard

Standards track:

Proposed Standard → Draft Standard → Standard

Åbne standarder = åbenhed, ikke garanti for sikkerhed

What is the Internet



Communication between humans - currently!

Based on TCP/IP

- best effort
- packet switching (IPv6 calls it packets, not datagram)
- *connection-oriented* TCP
- *connection-less* UDP

RFC-1958:


A good analogy for the development of the Internet is that of constantly renewing the individual streets and buildings of a city, rather than razing the city and rebuilding it. The architectural principles therefore aim to provide a framework for creating cooperation and standards, as a small "spanning set" of rules that generates a large, varied and evolving space of technology.

IP: Internet historically




- 1961 L. Kleinrock, MIT packet-switching teori
- 1962 J. C. R. Licklider, MIT - notes
- 1964 Paul Baran: On Distributed Communications
- 1969 ARPANET start 4 nodes
- 1971 14 nodes
- 1973 Work on Internet Protocols IP started
- 1973 Email is about 75% af ARPANET trafik
- 1974 TCP/IP: Cerf/Kahn: A protocol for Packet Network Interconnection
- 1983 EUUG → DKUUG/DIKU connection
- 1988 ca. 60.000 systemer på Internet The Morris Worm rammer ca. 10%
- 2002 Ialt ca. 130 millioner på Internet

Internet historically set - anno 1969



- Node 1: University of California Los Angeles
- Node 2: Stanford Research Institute
- Node 3: University of California Santa Barbara
- Node 4: University of Utah


What are Internet hosts



Cumulative RIR address assignments, per RIR

Source: IPv4 Address Report - 28-Jan-2019 <http://www.potaroo.net/tools/ipv4/>

Common Address Space



Internet is defined by the address space, one
Based on 32-bit addresses, example 10.0.0.1

IPv4 address



```
hlk@bigfoot:hlk$ ipconvert.pl 127.0.0.1
Adressen er: 127.0.0.1
Adressen er: 2130706433
hlk@bigfoot:hlk$ ping 2130706433
PING 2130706433 (127.0.0.1): 56 data bytes
64 bytes from 127.0.0.1: icmp_seq=0 ttl=64 time=0.135 ms
64 bytes from 127.0.0.1: icmp_seq=1 ttl=64 time=0.144 ms
```

IP-adresser typically written as decimal numbers with dots

dot notation: 10.1.2.3

IP-adresser as bits



IP-adresse: 127.0.0.1

Heltal: 2130706433

Binary: 11111110000000000000000000000001

IP-address converted to bits

Computers use bits

Internet ABC



Previously we used classes: A, B, C, D og E

This proved to be a bit inflexible:

- A-klasse has 16 million hosts
- B-klasse about 65.000 hosts
- C-klasse only 250 hosts

Most people asked for B-klasser - starting to run out!

D-klasse used for multicast

E-klasse reserved

See http://en.wikipedia.org/wiki/Classful_network

Stop saying C, say /24

CIDR Classless Inter-Domain Routing



Classful routing		Classless routing CIDR	
4 class C networks	Inherent subnet mask	Supernet	Subnet mask
192.0.8.0	255.255.255.0	192.0.8.0	255.255.252.0
192.0.9.0	255.255.255.0		252d=11111100b
192.0.10.0	255.255.255.0		
192.0.11.0	255.255.255.0	Base network/prefix	
		192.10.8.0/22	

Subnet mask originally inferred by the class

Started to allocate multiple C-class networks - save remaining B-class

Resulted in routing table explosion

A subnet mask today is a row of 1-bit

10.0.0.0/24 means the network 10.0.0.0 with subnet mask 255.255.255.0

Supernet, supernetting

RFC-1918 Private Networks



Der findes et antal adresserum som alle må benytte frit:

- 10.0.0.0 - 10.255.255.255 (10/8 prefix)
- 172.16.0.0 - 172.31.255.255 (172.16/12 prefix)
- 192.168.0.0 - 192.168.255.255 (192.168/16 prefix)

Address Allocation for Private Internets RFC-1918 adresserne!

NB: man må ikke sende pakker ud på internet med disse som afsender, giver ikke mening

The blocks 192.0.2.0/24 (TEST-NET-1), 198.51.100.0/24 (TEST-NET-2), and 203.0.113.0/24 (TEST-NET-3) are provided for use in documentation.

169.254.0.0/16 has been ear-marked as the IP range to use for end node auto-configuration when a DHCP server may not be found

OSI og Internet modellerne




OSI Reference
Model

Application
Presentation
Session
Transport
Network
Link
Physical

Internet protocol suite

Applications HTTP, SMTP, FTP, SNMP,	NFS
	XDR
	RPC
TCP UDP	
IPv4	IPv6 ICMPv6 ICMP
ARP RARP	
MAC	
Ethernet token-ring ATM ...	

Packets across the wire or wireless




Looking at data as a stream the packets are a pattern laid on top

Network technology defines the start and end of a frame, example Ethernet

From a lower level we receive a packet, example 1500-bytes from Ethernet driver

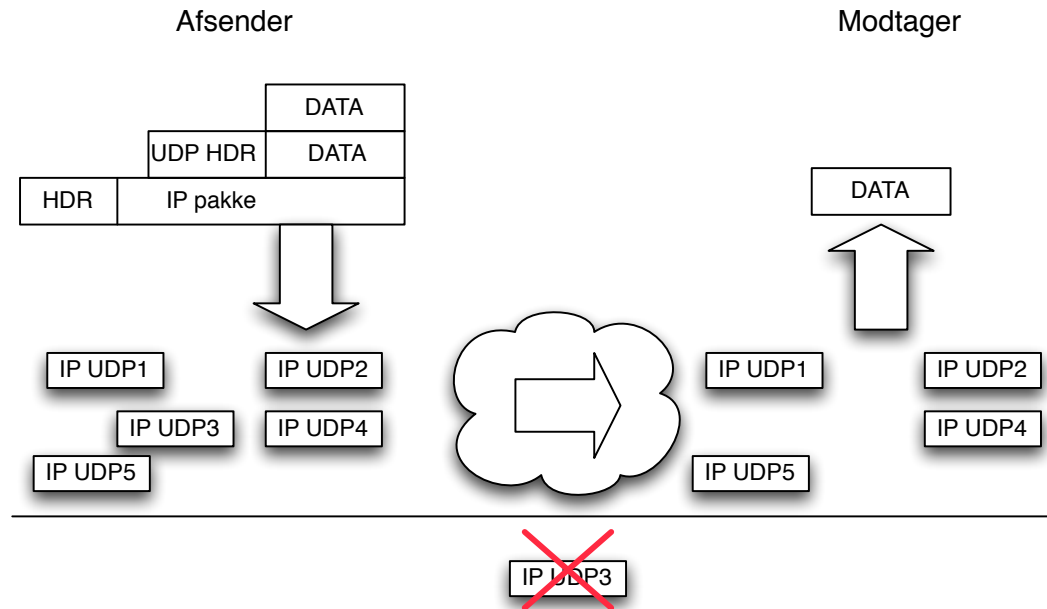
Operating system masks a lot of complexity

IPv4 pakken - header - RFC-791



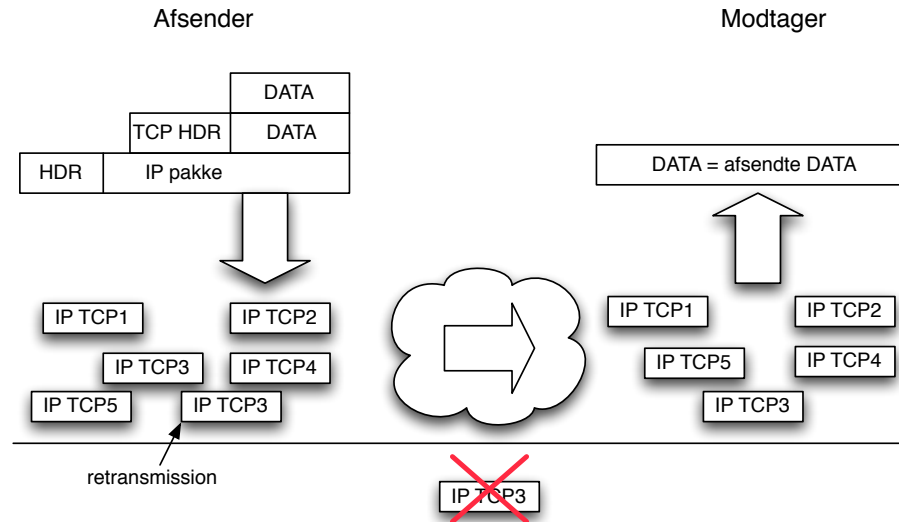
Example Internet Datagram Header

UDP User Datagram Protocol



Connection-less RFC-768, *connection-less*
Used for Domain Name Service (DNS)

TCP Transmission Control Protocol




Connection oriented RFC-791 September 1981, *connection-oriented*

Either data delivered in correct order, no data missing, checksum or an error is reported

Used for HTTP and others

TCP three way handshake



- Session setup is used in some protocols
- Other protocols like HTTP/2 can perform request in the first packet

Well-known port numbers



IANA maintains a list of magical numbers in TCP/IP

Lists of protocol numbers, port numbers etc.

A few notable examples:

- Port 25/tcp Simple Mail Transfer Protocol (SMTP)
- Port 53/udp and 53/tcp Domain Name System (DNS)
- Port 80/tcp Hyper Text Transfer Protocol (HTTP)
- Port 443/tcp HTTP over TLS/SSL (HTTPS)

Source: <http://www.iana.org>

Exercise: Communicate with HTTP



Try this - use netcat/ncat, available in Nmap package from [Nmap.org](https://nmap.org):

```
$ netcat www.zencurity.com 80
GET / HTTP/1.0
```

```
HTTP/1.1 200 OK
Server: nginx
Date: Sat, 01 Feb 2020 20:30:06 GMT
Content-Type: text/html
Content-Length: 0
Last-Modified: Thu, 04 Jan 2018 15:03:08 GMT
Connection: close
ETag: "5a4e422c-0"
Referrer-Policy: no-referrer
Accept-Ranges: bytes
...
```

Basic test tools TCP - Telnet and OpenSSL



Telnet used for terminal connections over TCP, but is clear-text

Telnet can be used for testing connections to many older protocols which uses text commands


- `telnet mail.kramse.dk 25` create connection to port 25/tcp
- `telnet www.kramse.dk 80` create connection to port 80/tcp

Encrypted connections often use TLS and can be tested using OpenSSL command line tool `openssl`

- `openssl s_client -host www.kramse.dk -port 443`
create connection to port 443/tcp with TLS
- `openssl s_client -host mail.kramse.dk -port 993`
create connection to port 993/tcp with TLS

Using OpenSSL in client-mode allows the use of the same commands like Telnet after connection

Book: Practical Packet Analysis (PPA)



Practical Packet Analysis, Using Wireshark to Solve Real-World Network Problems by Chris Sanders, 3rd Edition April 2017, 368 pp. ISBN-13: 978-1-59327-802-1

<https://nostarch.com/packetanalysis3>

Apache Tomcat



The Apache Tomcat® software is an open source implementation of the Java Servlet, JavaServer Pages, Java Expression Language and Java WebSocket technologies. The Java Servlet, JavaServer Pages, Java Expression Language and Java WebSocket specifications are developed under the Java Community Process.

- Allows the deployment of web applications J2EE
https://en.wikipedia.org/wiki/Java_Platform,_Enterprise_Edition
- Allows the use of Java security policies
- Contains the core functionality found in commercial packages
- <http://tomcat.apache.org/>

Java Apps, Tomcat, XML config



We will download Apache Tomcat, and perform the following:

- Download the software use version 9.0.30
I downloaded `apache-tomcat-9.0.30.tar.gz`, Windows users take the `.zip`
- Unpack and Run the software, see it works
- Work with Tomcat manager
- Check the configuration - which is in XML
- Change the configuration - make the software listen on all IPs, specific IP

Apache Tomcat Download



The screenshot shows a web browser window displaying the Apache Tomcat homepage. The browser's address bar shows the URL `tomcat.apache.org`. The page features the Apache Tomcat logo (a stylized orange cat) and the Apache Software Foundation logo. A search bar is located below the logo. The main content area is titled "Apache Tomcat" and contains several paragraphs of text describing the software. A sidebar on the left lists navigation links: "Home", "Taglibs", "Maven Plugin", "Download", and "Documentation". The "Download" section lists various versions and connectors. The "Documentation" section lists "Tomcat 9.0". A prominent announcement banner at the bottom of the main content area states "Tomcat 7.0.99 Released" with the date "2019-12-17".

Apache Tomcat®

Search... GO

APACHE EVENTS LEARN MORE

[Save the date!](#)

Apache Tomcat

- Home
- Taglibs
- Maven Plugin

Download

Which version?

- Tomcat 9
- Tomcat 8
- Tomcat 7
- Tomcat Connectors
- Tomcat Native
- Taglibs
- Archives

Documentation

- Tomcat 9.0

Apache Tomcat

The Apache Tomcat® software is an open source implementation of the Java Servlet, JavaServer Pages, Java Expression Language and Java WebSocket technologies. The Java Servlet, JavaServer Pages, Java Expression Language and Java WebSocket specifications are developed under the [Java Community Process](#).

The Apache Tomcat software is developed in an open and participatory environment and released under the [Apache License version 2](#). The Apache Tomcat project is intended to be a collaboration of the best-of-breed developers from around the world. We invite you to participate in this open development project. To learn more about getting involved, [click here](#).

Apache Tomcat software powers numerous large-scale, mission-critical web applications across a diverse range of industries and organizations. Some of these users and their stories are listed on the [PoweredBy](#) wiki page.

Apache Tomcat, Tomcat, Apache, the Apache feather, and the Apache Tomcat project logo are trademarks of the Apache Software Foundation.

Tomcat 7.0.99 Released 2019-12-17

The Apache Tomcat Project is proud to announce the release of version 7.0.99 of Apache Tomcat. This release contains a number of bug fixes and improvements compared to version 7.0.96.

Full details of these changes, and all the other changes, are available in the [Tomcat 7 changelog](#).

Unpack and run

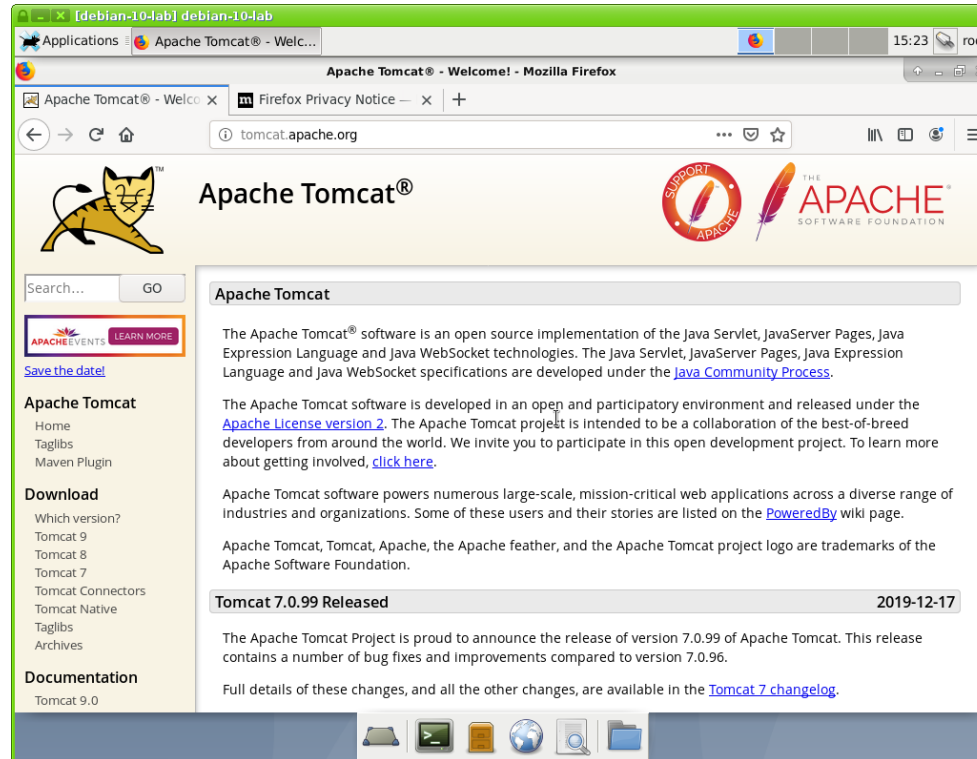


Debian instructions - shortened:

- `mkdir projects;cd projects -`
- `tar zxvf $HOME/Downloads/apache-tomcat-9.0.30.tar.gz -`
- `cd apache-tomcat-9.0.30 -`
- `sh bin/startup.sh -`

Windows - double click and unpack should produce the directory, and use the `startup.bat`

Tomcat Running



Now use a little time to look into the server, links, which ports are open - 8080/tcp and 8009/tcp

Tomcat Tasks



- Try going to `http://127.0.0.1:8080/manager/` - get 401 Unauthorized
Read and fix the problem
- Look into the XML file `conf/server.xml`
Change the port 8080 into something else, does port 80/tcp work?
- XML configuration files are very common in the Java and system integration space
- Tools often found in Java world, Ant, Maven, Tomcat, Spring
- Read <https://tomcat.apache.org/tomcat-9.0-doc/monitoring.html> about Java Management Extensions (JMX) https://en.wikipedia.org/wiki/Java_Management_Extensions
Note: references to RMI, JMS, SNMP, HTTP etc.

Git intro




Git (/ t/)[7] is a distributed version-control system for tracking changes in source code during software development.[8] It is designed for coordinating work among programmers, but it can be used to track changes in any set of files. Its goals include speed,[9] data integrity,[10] and support for distributed, non-linear workflows.[11]

Source: <https://en.wikipedia.org/wiki/Git>

- We will introduce Git and Github - commercial Git hosting company
<https://en.wikipedia.org/wiki/Git>
- Try creating a Git repository, remember to add a README
- Checkout the repository
- Edit a file
- add and commit it

Use the Github Hello World example: <https://guides.github.com/activities/hello-world/>

For Next Time



Think about the subjects from this time, write down questions

Check the plan for chapters to read in the books

Visit web sites and download papers if needed

Retry the exercises to get more confident using the tools